

Physics informed neural networks for problems involving nonlinear differential equations.

Project report

CS498DL, Spring 2021

1 Team details

Name	NetID	Department
Bhosale, Yashraj	bhosale2@illinois.edu	Mechanical Science and Engineering
Chan, Fan Kiat	fchan5@illinois.edu	Mechanical Science and Engineering
Parthasarathy, Tejaswin	tp5@illinois.edu	Mechanical Science and Engineering

2 Updates

For ease in readability and grading, updates to the report (from the first checkpoint) and new results [are marked in violet, like this sentence.](#)

3 Problem summary

We implement and investigate Physics-Informed Neural Networks, a deep-learning approach for solving non-linear partial differential equations. We demonstrate PINNs' ability to solve equations of increasing complexity, from linear, 1D problems to nonlinear, 2D problems, based only on the problem description and without the need of fine numerical discretization or problem-specific algorithms.



Figure 1: Solution of a Poisson equation generated by Physics-Informed Neural Networks within the shape *CS498DL*. Each alphanumeric character is prescribed with a unique boundary condition resulting in varied, rich solutions seen in the interior.

4 Introduction and background

In this report, we implement and investigate Physics-Informed Neural Networks (PINNs), a deep-learning approach for generating numerical solutions of non-linear partial differential equations (nPDEs). We are motivated by utilizing the recent rapid advances in data-driven techniques in machine learning (ML) and deep learning (DL) to solve nPDEs, as an alternative to conventional numerical approaches which are both mathematically and computationally intensive. Rather PINNs take minimal effort to set up, doesn't require domain knowledge (of PDEs) and is hence ideal as a black-box solver which once trained, can rapidly produce solutions to arbitrary numerical PDEs. In the rest of this section, we introduce the reader to PDEs and conventional numerical approaches to solve them, which in turn motivates the need for alternative methods such as PINNs.

4.1 Partial differential equations

Partial differential equations (PDE) are ubiquitous in science and engineering. Mathematically, a PDE imposes relations between various partial derivatives of a multi-variable function within a system of interest. These relations capture fundamental physical laws governing a system (e.g. Newton's equations of motion) such as those associated with sound, heat, diffusion, electrostatics, electrodynamics, fluid dynamics, elasticity, general relativity, and quantum mechanics [1].

Such PDEs are broadly categorized as initial-value problems (IVPs) and boundary-value problems (BVPs). Roughly speaking, IVPs involve *evolution* of a function as time progresses, from some initial condition, such as a trajectory function of a ball dropped from a height. In contrast, BVPs involve finding a static solution *satisfying constraints* on given boundaries, such as finding the temperature distribution of a room with a constant temperature heater embedded in one wall.

Many PDEs of practical interest (such as those governing fluid-dynamics) demonstrate behavior of both IVPs and BVPs. The solution to such PDEs are crucial in revealing physics to enable new applications (such as targeted therapeutic drug-delivery within our fluidic blood vessels, for curing diseases) or optimizing design of existing applications (such as improving models of aero-planes for efficient transport) or for predictive forecasting (such as predicting weather). Unfortunately, these PDEs are usually non-linear (i.e. they belong to the category of non-linear partial differential equations, nPDEs), and we do not *a priori* know their closed-form mathematical solutions.

4.2 Numerical solution to partial differential equations

To solve these nPDEs then, we turn to numerical solutions. The traditional approach here is to numerically *discretize* the nPDE over a large grid of floating-point variables in space-time. We then numerically solve, via computer simulations, the PDE locally at every point in the grid. Finally, we reconstruct the PDE solution using these local discretizations. This field has a rich tradition over the past ~70 years [2] and multiple techniques for efficiently solving IVPs and BVPs exist. As a result it has been quite successful in achieving its goals, especially with the recent advent of modern, parallel supercomputers which provide capabilities for accelerated simulations.

Despite its varied successes, this field is not without challenges. To achieve satisfactory results, the mesh spacing of the grid needs to be smaller than the smallest feature size of the target solutions. This often is not feasible because of the curse of dimensionality: achieving $10\times$ higher resolution usually requires $10^4\times$ more compute, because the grid must be scaled in four dimensions—three spatial and one temporal. Then, to satisfactorily solve these numerical equations we need faster and larger simulations. But in recent years, Moore's (predictive) law has been slowing [3] and we are approaching the limits of Dennard scaling [4], which restricts design of general-purpose microprocessors, which consequently stifles simulation speeds even with parallelism. [Another challenge](#)

relates to the wide variety of numerical algorithms needed. In fact, different PDEs require vastly different numerical strategies to ensure physical fidelity of *discrete* solutions in the computer. Developing the theory for such strategies is usually mathematically cumbersome and time intensive, especially when new PDEs need to be modeled.

4.3 PINNs

An alternative approach to solving PDEs numerically, aiming to partially offset the aforementioned challenges, involves utilizing data-driven techniques popularized by the advent of machine learning (ML) [5]. Such approaches efficiently utilize data sampled from systems of interest to learn the representation of underlying nPDEs. They also take advantage of recent algorithmic breakthroughs in ML along with fast hardware (such as TPUs) and software (such as frameworks) optimized for these applications.

An attractive candidate among these approaches is the Physics-Informed Neural Network (PINN). These are models trained to solve self-supervised learning tasks adhering to given laws of physics [6] described by general nPDEs. As such, only the initial and boundary data are needed to recover solutions to arbitrary nPDEs—the network generates incorrect solutions while training and uses these solutions to learn, and converge to the correct result. *Additionally, they are easy to setup and efficiently run on modern hardware architectures, thanks to automatic differentiation and abstracted kernel (GPU) implementations in recent, popular deep learning software frameworks [7].* These features make it an attractive candidate for use as a black-box multi-physics solver across engineering domains. Finally, recent works on PINN have shown promise in *instantly* predicting results for IVP governed by nPDEs across scientific domains, for example in incompressible fluid dynamics [8, 9], reaction–diffusion systems in computational chemistry [6], and the propagation of nonlinear waves in quantum mechanics. PINNs can also generate solutions for BVPs in complex domains, such as those found in continuum solid mechanics.

4.4 Deliverables

Having motivated the need for ML approaches such as PINNs for solving PDEs,

1. We propose to implement the algorithm for such PINNs and validate it against a time-independent BVPs with known (numerical) solutions (such as Laplace or Poisson equation), as well as against time-dependent IVPs (such as the Burgers equation).
2. Finally, if time permits, we will extend the method for solving problems of higher complexity. In particular, we are interested in solving problems in the field of fluid mechanics [8, 9], with focus on problems where no known analytical solutions are available and direct numerical simulation approaches are computationally expensive.

5 Results

We demonstrate our results by considering PINN solutions to increasingly complex PDEs in a step-by-step fashion. Particularly, we start by describing the implementation of such PINNs. Then we consider a sequence of PINN solutions to 1D BVPs (Poisson, Helmholtz), and build up our PINN machinery which we utilize to generate solutions to more complex spatio-temporal nPDEs (Burgers, 2D Poisson).

5.1 Implementation details

Here we consider BVPs as parametrized PDEs of the form

$$\mathcal{P}[u; \lambda] = 0, x \in [0, 1]$$

where x denotes the solution domain, $u : x \mapsto \mathbb{R}$ denotes the latent solution embedded in our deep network, and $\mathcal{P}[\cdot, \lambda]$ is the operator defining the PDE, parametrized by our PINN weights (and biases) λ . These parameters λ are learnt over a self-supervised training process where we only provide the PDE to be satisfied and boundary data. We now describe our network architecture, loss functions and training protocols used to implement PINNs.

5.1.1 Network architecture

We design our PINNs to take input points x_i , $i = 1, \dots, N_{\mathcal{P}}$ sampled $\in [0, 1]$ as a vector $\in \mathbb{R}^{N_{\mathcal{P}}}$. The total number of input points $N_{\mathcal{P}}$ is chosen based on the desired resolution. The network then generates potential solutions $u_i \in \mathbb{R}^{N_{\mathcal{P}}}$ as outputs, sampled at the same points.

Throughout this report our PINN utilizes a multi-layer perceptron, comprising of 2 hidden, fully-connected layers of 50 neurons each with a hyperbolic tangent (tanh) activation function. Our choice of tanh activation is necessary to accelerate learning by allowing back-propagation of non-zero gradients, thus bypassing dead ReLU problems. Other complex activation functions [10] are possible—here we skip these to retain simplicity in our networks.

Finally, after these hidden layers, we have a fully-connected 50-neuron output layer bereft of any activation. The lack of final layer activation is motivated by our need to capture arbitrary PDE and boundary-value dependent scales and shifts of PDE solutions. If we enabled a similar tanh activation here, it puts the total network output in the limited range $[-1, 1]$, which curtails the representation power of our network and consequently destroys any potential for generalization to arbitrary PDEs. With this fixed model, we proceed to train our network.

5.1.2 Loss function

During training, we minimize the loss of sum of mean squared errors associated with the PDE \mathcal{P} and boundary data u_b . That is the total loss MSE is defined as

$$\text{MSE} := \lambda_{\mathcal{P}} \text{MSE}_{\mathcal{P}} + \text{MSE}_b$$

where the *PDE loss* $\text{MSE}_{\mathcal{P}}$ is defined as

$$\text{MSE}_{\mathcal{P}} := \mathcal{L}_2 [\mathcal{P}[u(x_i)]]$$

where $\mathcal{L}_2[\cdot]$ is the discrete L2 norm defined as

$$\mathcal{L}_2[f] := \frac{1}{N_{\mathcal{P}}} \sum_{i=1}^{N_{\mathcal{P}}} |f(x_i)|^2$$

The PDE loss above enforces adherence of latent solution u to \mathcal{P} across points sampled within the domain x_i , $i = 1, \dots, N_{\mathcal{P}}$. The second loss MSE_b enforces adherence of u to boundary values

and is called the *boundary loss*, defined as

$$\text{MSE}_b := \frac{1}{N_b} \sum_{j=1}^{N_b} \left| u(x_j^b) - u_b(x_j^b) \right|^2$$

where N_b denotes the number of boundary points and the sum here is over all boundary points x_j^b , $j = 1, \dots, N_b$.

To adjust the relative importance the network provides to conformation to the PDE in the interior points versus the boundary points, we use the hyper-parameter $\lambda_{\mathcal{P}}$, which we set to 1 in our experiments, unless noted otherwise.

5.1.3 Training protocol

One attractive feature of PINNs for generating PDE solutions is that its training is not data intensive. Indeed, PINNs are self-supervised and works in sparse-data limits—it only needs the PDE structure and boundary information! Hence, in this project, we do not train it over huge datasets, as a result of which training time is manageable, taking only a [couple of minutes across all presented cases](#), on both our local machines and Google Colaboratory servers.

Training for all architectures presented here is done with the Adam optimizer [11] with a training rate of 2.5×10^{-3} , for 15000 epochs, with other parameters set to their recommended default.

Once the training is completed, we then need to test fitness of generated solutions u . For linear PDEs, we have analytical solutions \hat{u} , which depends on the problem and boundary conditions, against which we compare u . This \hat{u} is specified as and when needed, from Section 5.2 onwards. It is also sampled along the network points x_i , $i = 1, \dots, N_{\mathcal{P}}$. This allows us to define a network error $e := \sqrt{\mathcal{L}_2[u - \hat{u}]}$, which we use to gauge solution fitness—smaller errors are better.

In contrast, for nonlinear PDEs, we lack analytical solutions. Here we then adopt two metrics to gauge solution fitness:

- We measure the residual $\sqrt{\mathcal{L}_2[\mathcal{P}[u]]}$. This indicates how well u satisfies the PDE. Smaller residuals are better.
- We obtain numerical solutions u_h , using finite-differences, against which we compare u , to produce discrete error e_h defined as $e_h := \sqrt{\mathcal{L}_2[u - u_h]}$. Once again we use this discrete error e_h to gauge solution fitness—smaller errors are better.

A fit solution is marked by small magnitudes across both these metrics.

This concludes our section on model choice and training. We now use this model to investigate performance of PINNs for generating solutions to PDEs of increasing complexity.

5.2 PINN solutions of 1D Poisson equation

We start by investigating PINN solutions of simple linear PDEs with established theoretical results against which we compare. Here, we consider solutions $u : x \mapsto \mathbb{R}, x \in [0, 1]$ to the 1D Poisson equation, described by

$$\begin{aligned} \mathcal{P}[u] &:= \nabla^2 u - f = 0 \\ &:= \frac{\partial^2 u}{\partial x^2} - f(x) = 0 \end{aligned} \tag{1}$$

coupled with suitable boundary conditions. We start with the Poisson equation as it is the workhorse underlying virtually all scientific and engineering domains. As a more concrete example, it is useful in heat-transfer applications such as thermal regulation of indoor spaces. Here

the Poisson Eq. (1) governs the distribution of temperature (modeled using u) within a room with heat sources/sinks (modeled using f) with, say, insulated walls (modeled using boundary conditions). Here we do not model such applications in full complexity, rather we focus on a toy problem capturing the essence of generating solutions to the 1D Poisson Eq. (1).

In this equation, we first demonstrate our ability to account for varied boundary conditions in our PINN architecture for a fixed function f :

$$f(x) = 10 (\sin(\pi x) + 2 \sin(2\pi x) + 3 \sin(3\pi x) + 10 \sin(4\pi x) + 10 \sin(5\pi x))$$

This sinusoidal forcing is chosen for producing relatively complex solutions with multiple bumps and troughs, despite its simplicity in description. With this $f(x)$ we then compare our PINN results against known analytical solutions.

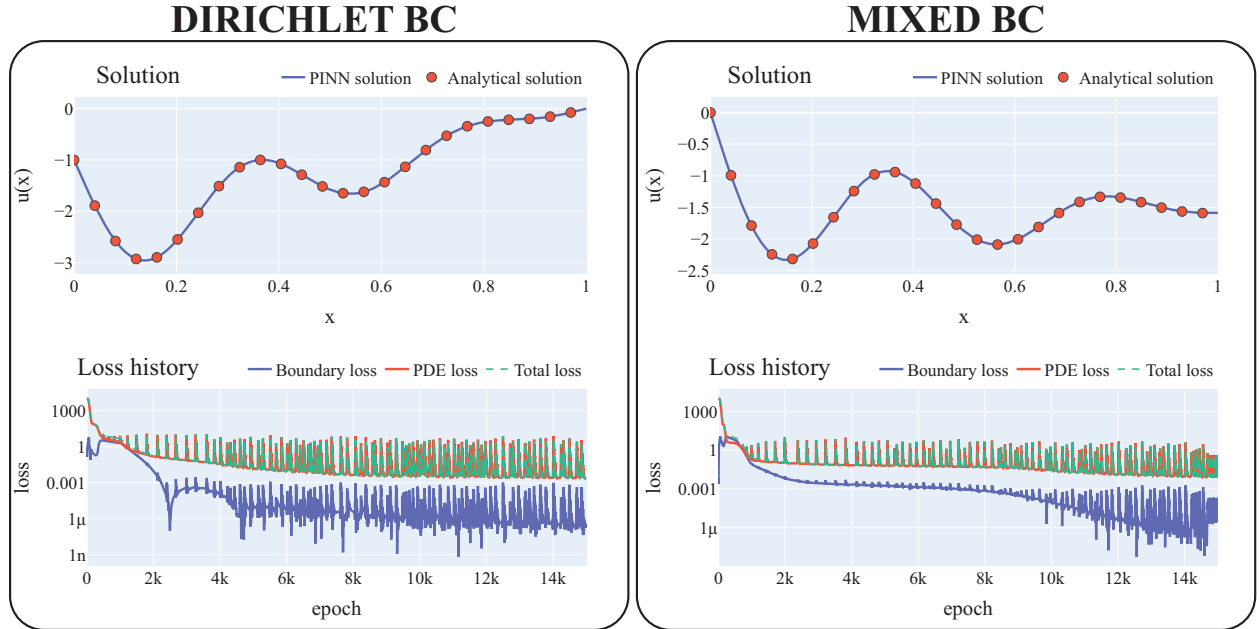


Figure 2: PINN results (solution and loss histories) for the 1D Poisson equation shown for (left) Dirichlet boundary conditions and (right) mixed Dirichlet–Neumann boundary condition. The PINN solutions are depicted by solid, blue lines while scatter points represent reference analytical solutions.

5.2.1 Dirichlet boundaries

In this regard, we first solve Eq. (1) with Dirichlet (or) constant-value boundary conditions

$$u_b(0) = -1; \quad u_b(1) = 0$$

For this case, our analytical solution is

$$\hat{u}(x) = -1 + x - \frac{10}{\pi^2} \left(\sin(\pi x) + \frac{1}{2} \sin(2\pi x) + \frac{1}{3} \sin(3\pi x) + \frac{5}{8} \sin(4\pi x) + \frac{2}{5} \sin(5\pi x) \right)$$

We train our PINN with $N_{\mathcal{P}} = 100, N_b = 2$ and recover the solutions shown in the left column of Fig. 2, along with its loss histories. We see that our PINN solution agrees with the analytical one, with $e = 4.6 \times 10^{-4}$. The loss history is dominated by $\text{MSE}_{\mathcal{P}}$ and steadily decreases with increasing epochs.

5.2.2 Mixed Dirichlet–Neumann boundaries

Next we demonstrate solutions of Eq. (1) with Neumann (or) constant-slope boundary condition on the right boundary and retain Dirichlet conditions on the left boundary.

$$u(0)_b = 0; \quad \frac{\partial u_b}{\partial x}(1) = 0$$

For this case, our analytical solution is

$$\begin{aligned} \hat{u}(x) = & \frac{10}{\pi} \left(\cos(\pi) + \cos(2\pi) + \cos(3\pi) + \frac{5}{2} \cos(4\pi) + 2 \cos(5\pi) \right) x \\ & - \frac{10}{\pi^2} \left(\sin(\pi x) + \frac{1}{2} \sin(2\pi x) + \frac{1}{3} \sin(3\pi x) + \frac{5}{8} \sin(4\pi x) + \frac{2}{5} \sin(5\pi x) \right) \end{aligned} \quad (2)$$

We again train our PINN with $N_{\mathcal{P}} = 100, N_b = 2$ and recover the solutions shown in the right column of Fig. 2, along with its loss histories. We see that our PINN solution agrees with the analytical one, with $e = 2.7 \times 10^{-4}$. The loss history is dominated by $\text{MSE}_{\mathcal{P}}$ and steadily decreases with increasing epochs.

5.3 PINN solutions of 1D Helmholtz equation

Next, we consider solutions to more complex linear problems. Here, we consider solutions $u : x \mapsto \mathbb{R}, x \in [0, 1]$ to the 1D Helmholtz equation, described by

$$\begin{aligned} \mathcal{P}[u] &:= (\nabla^2 + k^2) u - f = 0 \\ &:= \frac{\partial^2 u}{\partial x^2} + k^2 u - f(x) = 0 \end{aligned} \quad (3)$$

coupled with suitable boundary conditions. Here, $k \in \mathbb{R}$ denotes a constant wave number and introduces additional complexity to Eq. (3) when compared to Eq. (1). Helmholtz equations are typical in scenarios involving waves (for e.g. in acoustics and electromagnetic radiation) and subsequent applications (for e.g. in noise-cancelling headphones). Here it governs the amplitude of sound (modeled using u) within a room with noise sources/sinks (modeled using f) with, say, perfectly sound-proof walls (modeled using boundary conditions).

While we demonstrated our ability to handle different boundary conditions in Section 5.2, here we demonstrate our ability to account for arbitrary forcing functions $f(x)$ in our PINN architecture. We achieve this for fixed, Dirichlet boundary conditions

$$u_b(0) = 0; \quad u_b(1) = 0$$

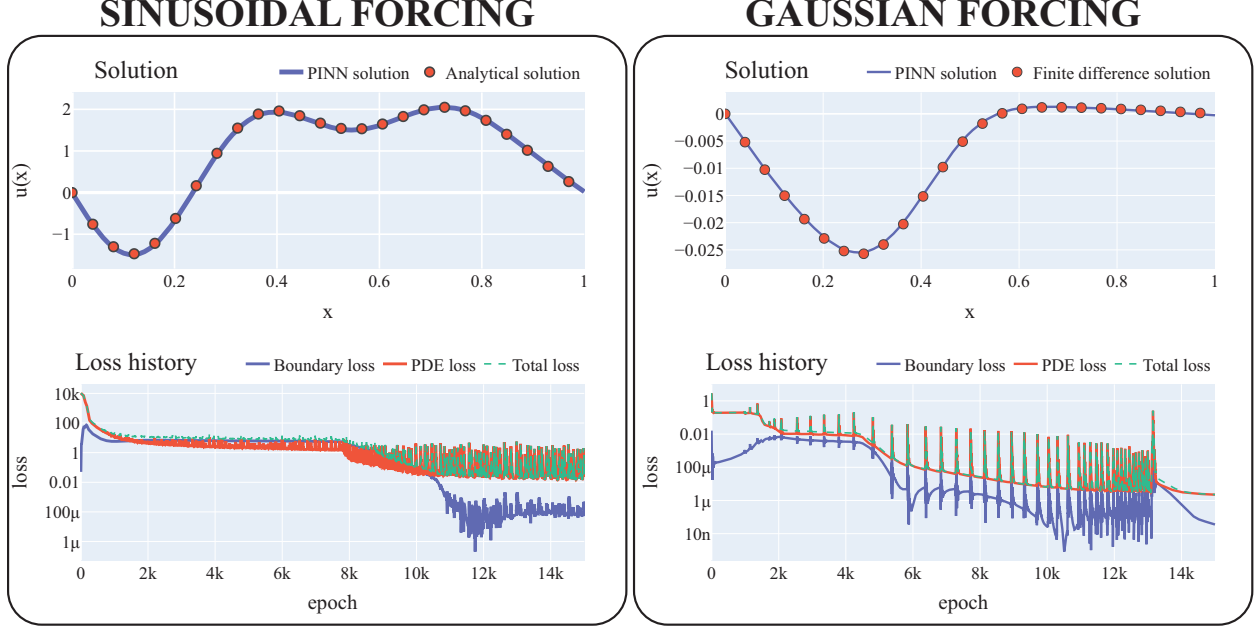


Figure 3: PINN results (solution and loss histories) for the 1D Helmholtz equation shown for (left) sinusoidal forcing and (right) gaussian forcing. The PINN solutions are depicted by solid, blue lines while scatter points represent reference analytical/finite-difference solutions.

5.3.1 Sinusoidal forcing

In this regard, we first solve Eq. (3) for $k = 4$ with the same forcing equation as in Section 5.2, reproduced below for convenience.

$$f(x) = 10 (\sin(\pi x) + 2 \sin(2\pi x) + 3 \sin(3\pi x) + 10 \sin(4\pi x) + 10 \sin(5\pi x))$$

For this case, our analytical solution is

$$\hat{u}(x) = 10 \left(\frac{1}{4^2 - \pi^2} \sin(\pi x) + \frac{2}{4^2 - (2\pi)^2} \sin(2\pi x) + \frac{3}{4^2 - (3\pi)^2} \sin(3\pi x) + \frac{10}{4^2 - (4\pi)^2} \sin(4\pi x) + \frac{10}{4^2 - (5\pi)^2} \sin(5\pi x) \right) \quad (4)$$

We train our PINN with $N_{\mathcal{P}} = 100, N_b = 2$ and recover the solutions shown in the left column of Fig. 3, along with its loss histories. We see that our PINN solution agrees with the analytical one, with $e = 2 \times 10^{-2}$. Regarding training history, we see that the loss is dominated by MSE_b at early stages of training, and later dominated by $\text{MSE}_{\mathcal{P}}$. In other words, the PINN has difficulty adjusting to boundary conditions in the first 10k epochs, beyond which it finds the right set of solutions, marked by a decrease in boundary loss. For the PDE loss, we observe a similar decrease around 8k epochs but this time the behavior is comparatively less pronounced.

5.3.2 Gaussian forcing

Next we demonstrate solutions of Eq. (3) with a non-sinusoidal forcing demonstrated below, consisting of two Gaussian bumps

$$f(x) = \exp \left[- \left(\frac{(x - 0.3)}{0.1} \right)^2 \right] - \exp \left[- \left(\frac{(x - 0.5)}{0.1} \right)^2 \right]$$

In this case analytical solutions are not easy to obtain, so we turn to numerical solution u_h based on finite-differences. We again train our PINN with $N_{\mathcal{P}} = 100, N_b = 2$ and recover the solutions shown in the right column of Fig. 3, along with its loss histories. We see that our PINN solution agrees with the reference solution, with an $e_h = 2.1 \times 10^{-4}$. The loss history is dominated by $\text{MSE}_{\mathcal{P}}$ and steadily decreases with increasing epochs.

5.4 PINN solutions of 1D stationary viscous Burgers equation

Now, towards realizing our goals of solving temporally evolving non-linear problems, we introduce a non-linearity in the governing PDE. Here, we build up on results presented in Section 5.2 and Section 5.3, by adding a square non-linearity to the Poisson equation, similar in structure to the Helmholtz equation i.e. an additional term depending only on u^2 . The simplest PDE that satisfies these conditions, while being practically relevant, is the stationary viscous Burgers equation. We consider its solutions $u : x \mapsto \mathbb{R}, x \in [0, 1]$, described by

$$\mathcal{P}[u] := \frac{1}{Pe} \frac{\partial^2 u}{\partial x^2} + \frac{1}{2} \frac{\partial (u^2)}{\partial x} - f(x) = 0 \quad (5)$$

coupled with suitable boundary conditions. This stationary viscous Burgers equation is useful in predicting steady 1D velocity profiles over long plates in fluid-dynamics applications. One such application occurs in boundary-layer flow past an airplane, where we need to determine the velocity u to estimate drag forces (and fuel consumption), given the influence of atmosphere f . Here the Péclet number Pe is a fixed parameter $\in \mathbb{R}$ that moderates the effect of non-linearity in the physics of the problem. Higher Péclet numbers usually display interesting, non-linear behaviors.

We have already seen the ability of PINNs to handle different boundary conditions and arbitrary forcing. Then, in this section we focus on capturing non-linear effects moderated by the Péclet number Pe for a fixed forcing $f(x)$ and boundary conditions u_b based on realistic physics. Accordingly, we set

$$f(x) = 1$$

which mimics a constant atmospheric pressure gradient and set

$$u_b(0) = 0; \quad u_b(1) = 0$$

to enforces absence of velocities on the walls at $x = 0, 1$. To compare against our PINN solutions in these settings, we once again rely on numerical solutions u_h based on finite-differences.

5.4.1 $Pe = 5$

We begin by investigating the capability of our PINN to solve Eq. (5) for $Pe = 5$. Here, we expect non-linearities to be not too significant, and hence our network should behave similar to Section 5.2 and Section 5.3. Upon training our PINN with $N_{\mathcal{P}} = 100, N_b = 2$ we see the expected behavior

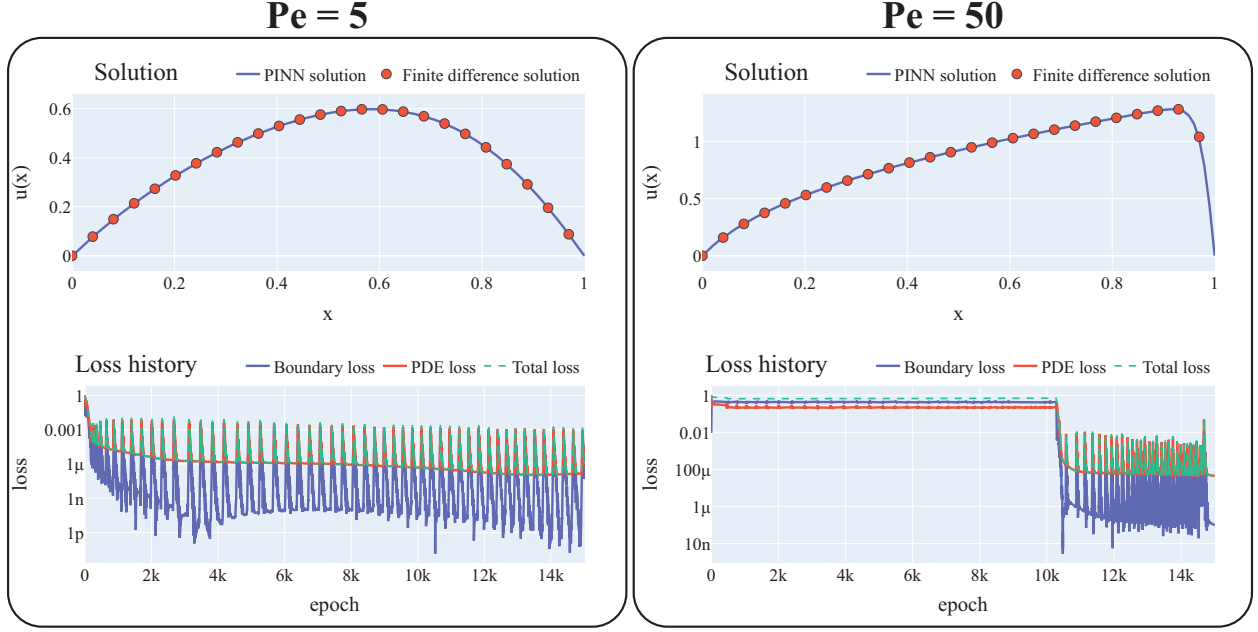


Figure 4: PINN results (solution and loss histories) for the 1D stationary viscous Burgers equation shown for a case with (left) $Pe = 5$ and (right) $Pe = 50$. The PINN solutions are depicted by solid, blue lines while scatter points represent reference finite-difference solutions.

and recover solutions shown in the left column of Fig. 4, along with its loss histories. We see that our PINN solution agrees with the reference solution, with $e_h = 6.4 \times 10^{-4}$. The loss history is dominated by $MSE_{\mathcal{P}}$ and steadily decreases with increasing epochs.

5.4.2 $Pe = 50$

Next, we investigate the case with $Pe = 50$, where effects of non-linearities are expected to be significant. Here we expect a relatively smooth solution in the bulk of the domain, with a rapid adjustment to the boundary condition at $x = 1$. Upon training our PINN with $N_{\mathcal{P}} = 100$, $N_b = 2$ we see this expected behavior and recover solutions shown in the right column of Fig. 4, along with its loss histories. We see that our PINN solution agrees with the reference solution, with an $e_h = 4.2 \times 10^{-3}$, indicating our ability to capture non-linearities. Regarding its training, the PINN barely learns till around 10k epochs, as the loss is almost constant. Then, after this critical epoch number, it discovers the solution manifold and the loss rapidly decreases. Beyond this point, the loss history is dominated by $MSE_{\mathcal{P}}$ and maintains a near-steady value. Such network behavior is unique to this case, and more investigations are necessary to uncover the network dynamics before and after the critical epoch.

5.5 PINN solutions of viscous Burgers equation

We then increase the complexity of PDE from our previous section on stationary viscous Burgers Eq. (5) by adding in a temporal evolution term. This results in the time evolving, nonlinear viscous

Burgers equation whose solutions $u : (y, t) \mapsto \mathbb{R}$, $y \in [0, 1]$, $t \in [0, 1]$, are described by

$$\mathcal{P}[u] := \frac{\partial u}{\partial t} + \frac{1}{2} \frac{\partial (u^2)}{\partial y} - \frac{1}{Pe} \frac{\partial^2 u}{\partial y^2} = 0 \quad (6)$$

coupled with suitable boundary conditions and an additional initial condition. Similar to the stationary Burgers equation, viscous Burgers equation is useful in modeling fluid-dynamics applications. Additionally, it finds use in nonlinear acoustics and even traffic-flow problems [12]! Indeed, typical solution to Eq. (6) form *shock waves* which are regions where the solution *folds* over itself due to non-linearities (such shocks are depicted in our results below). As a result of such folding, the solution changes behavior dramatically within a short span of space-time which then makes the resolution of such shocks non-trivial. Here we challenge our PINNs to reproduce the shock behavior characteristic of Eq. (6).

We have already seen the ability of PINNs to capture non-linear effects moderated by the Péclet number Pe for realistic conditions in Section 5.4, in the context of viscous Burgers equation. There, the solution was driven by an external forcing f . Here however, this forcing is absent. The solution evolution is then purely driven by time t , whose effects we access by modifying the solution profile at the initial time $t = 0$ for Dirichlet boundary conditions

$$u_b(0, t) = 0; \quad u_b(1, t) = 0$$

to enforces absence of velocities on the walls at $y = 0, 1$. To compare against our PINN solutions in these settings, we once again rely on numerical solutions u_h based on finite-differences.

5.5.1 Shock waves

We begin by investigating the capability of our PINN to capture non-linear shocks in Eq. (6). To instigate these shocks, we choose the following initial solution profile

$$u(y, 0) = \sin 2\pi y$$

We train our PINN with $N_{\mathcal{P}} = 10000$ sampled on a 100×100 uniform grid of space-time. We set $N_b = 300$ sampled uniformly on the three boundaries—two for enforcing the boundary condition, and one for enforcing the initial condition—each having 100 points. We showcase the results obtained in the left column of Fig. 5. First, we show the solution generated by the PINN on the top row. Here we see the characteristic nature of shocks, with the red and blue regions come together and collapse at the middle of the domain $y = 0.5$. Next, we show the spatial distribution of the relative errors (as percentages) obtained upon comparison with the numerical solution. As seen from this plot, the maximum errors are bounded $< 3\%$, and close to the shock location. Averaging these errors across the domain, we obtain $e = 6 \times 10^{-3}$. This confirms PINN’s ability to capture PDE solutions in spatio-temporal problems. Next, we present the PINN loss histories during training. Here, the loss history is dominated by $\text{MSE}_{\mathcal{P}}$ which decreases with increasing epochs. This loss reflects the localized PDE errors seen in the center of our space-time domain.

5.5.2 Rarefaction waves

Next, we investigate the formation of non-linear rarefaction waves with our PINN in Eq. (6). To setup these rarefaction waves, we choose the following initial solution profile

$$u(y, 0) = -\sin 2\pi y$$

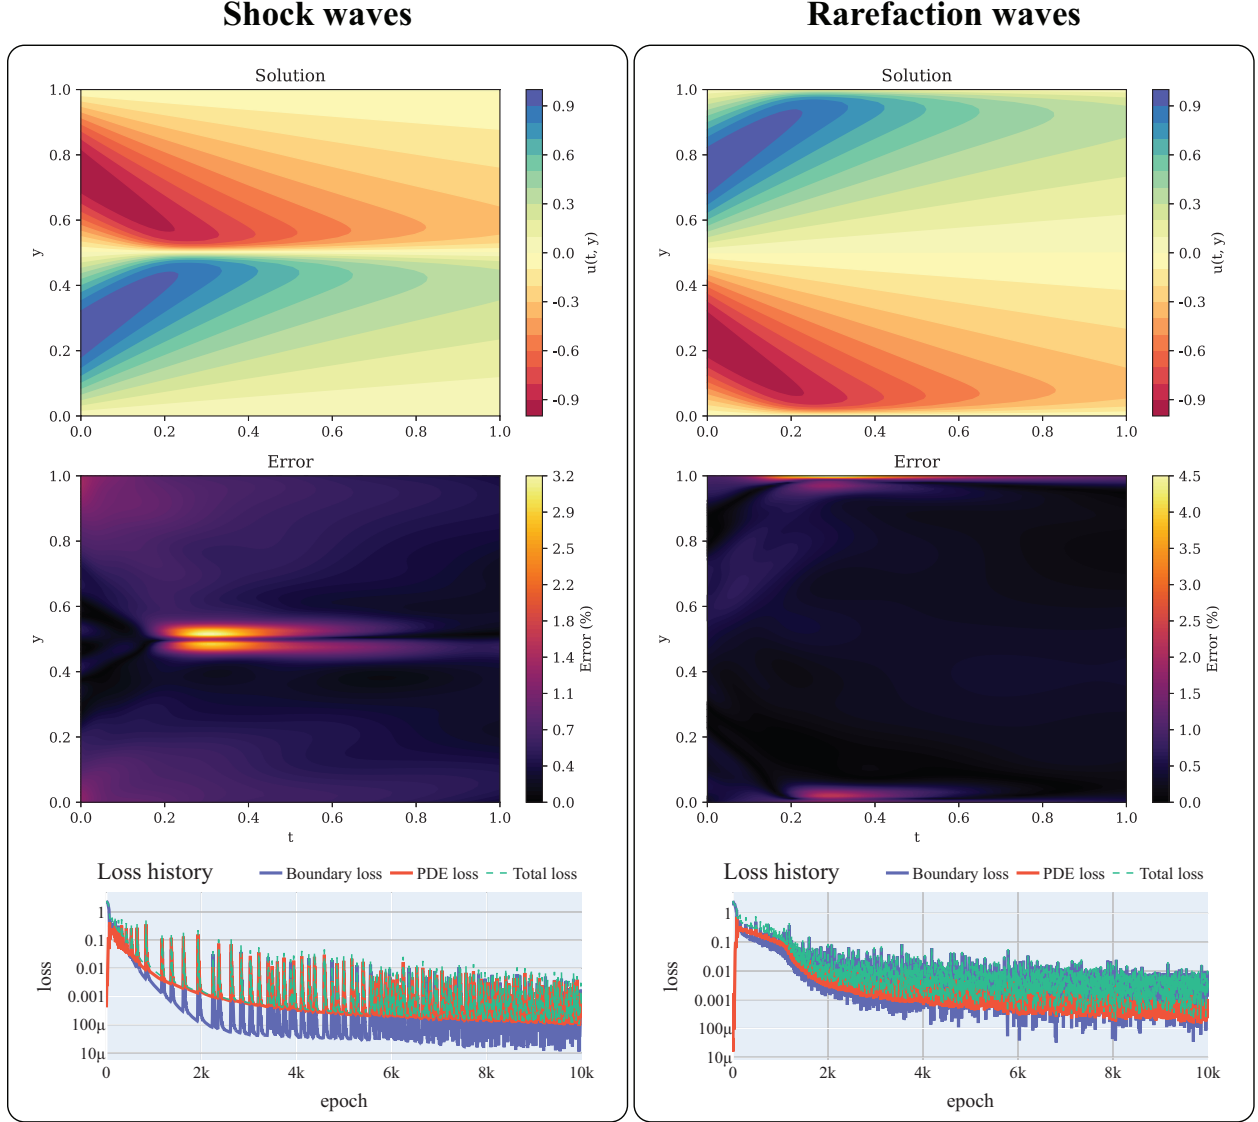


Figure 5: PINN results (solution and loss histories) for the viscous Burgers equation (varying with space–time) shown for different initial profiles that develop a (left) shock and (right) rarefaction waves. We first show the PINN solutions, followed by error plots (which we obtain by comparing to reference numerical solutions). Finally, loss histories are shown.

Once more, we train our PINN with $N_{\mathcal{P}} = 10000$ sampled on a 100×100 uniform grid of space-time. We set $N_b = 300$ sampled uniformly on the three boundaries—two for enforcing the boundary condition, and one for enforcing the initial condition—each having 100 points. We showcase the results obtained in the right column of Fig. 5. First, we show the solution generated by the PINN on the top row. We see the characteristic nature of rarefaction waves, where the red and blue regions which are initially close, diverge away from each other and fan out from the middle of the domain at $y = 0.5$. We then show the spatial distribution of the relative errors (as percentages) obtained upon comparison with the numerical solution. Once again the maximum errors are bounded $< 5\%$, but this time they are concentrated close to the domain boundaries. Averaging these errors across the domain, we obtain $e = 4.8 \times 10^{-3}$. This provides additional evidence that PINNs can viably capture non-linear spatio-temporal PDE solutions. Next, we present the PINN loss histories during training. Here, the loss history is dominated by MSE_b which decreases with increasing epochs. This loss reflects the localized errors seen near the boundary of our space-time domain.

5.6 PINN solutions of 2D Laplace equation

Finally we utilize the machinery of Eq. (6) to demonstrate solutions of the Laplace equation in fully two-dimensional settings. The solutions $u : \mathbf{x} := (x, y) \mapsto \mathbb{R}$, $\mathbf{x} \in \mathcal{D}$ for some domain \mathcal{D} are described by

$$\begin{aligned} \mathcal{P}[u] &:= \nabla^2 u = 0 \\ &:= \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0 \end{aligned} \tag{7}$$

coupled with suitable boundary conditions. The Laplace equation is a special case of the Poisson equation (see Section 5.2) with $f(\mathbf{x}) = 0$, and hence has similar applications.

While we have demonstrated the ability of PINNs to capture this PDE in Section 5.2, there we were limited to a rectangular domain in one-dimension. Here we crucially extend it to higher-dimensions (2D) in complex geometries of interest in design and engineering scenarios. In these settings, the boundary conditions can be spatially varying and hence the resulting solution is more challenging to capture.

5.6.1 Rectangular domain

We begin by testing our PINN architecture to solve Eq. (7) in a rectangular domain. Here we fix Dirichlet boundary condition in the x -direction:

$$u_b(0, y) = 0; \quad u_b(1, y) = 0$$

while in the y -direction we have a mixed boundary condition, similar to Section 5.2:

$$u_b(x, 0) = 0; \quad \frac{\partial u_b}{\partial x}(x, 1) = \pi (\exp(\pi) + \exp(-\pi)) \sin(\pi x)$$

This case lends itself to a straightforward analytical solution:

$$\hat{u}(x, y) = (\exp(\pi y) - \exp(-\pi y)) \sin(\pi x)$$

We train our PINN with $N_{\mathcal{P}} = 10000$ sampled on a 100×100 uniform grid within the domain. We set $N_b = 400$ once again sampled uniformly on the four boundaries having 100 points each. We

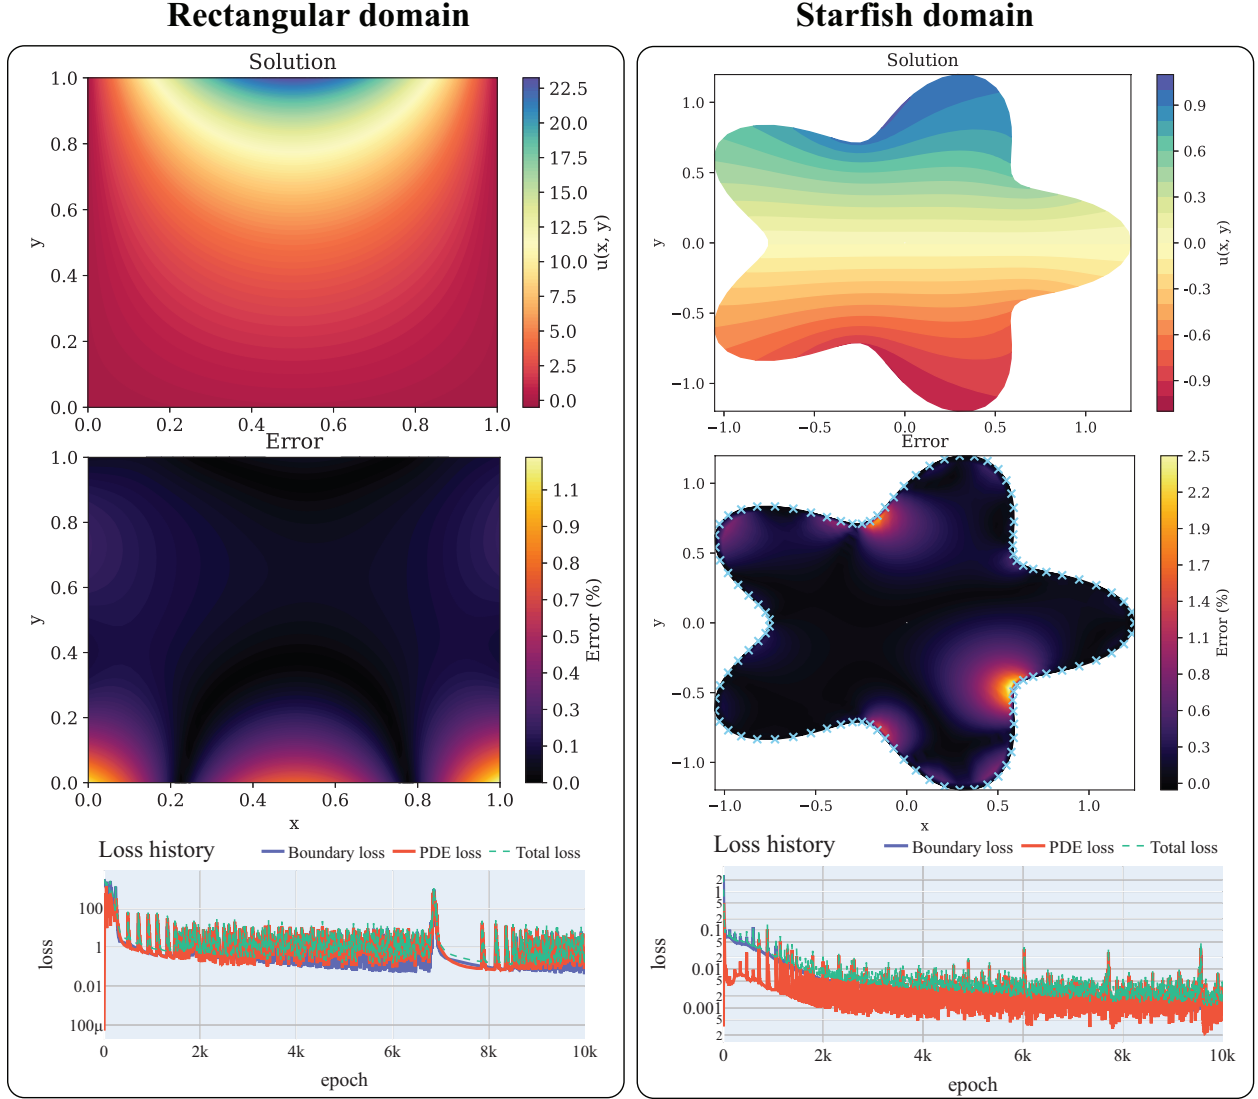


Figure 6: PINN results (solution and loss histories) for the 2D Laplace equation shown for a case in a (left) rectangular domain and (right) starfish domain. We first show the PINN solutions, followed by error plots (which we obtain by comparing to reference numerical solutions). Finally, loss histories are shown.

showcase the results obtained in the left column of Fig. 6. Here we first show the solution generated by the PINN on the first row. Next, we show the spatial distribution of the relative errors (as percentages) obtained upon comparison with the analytical solution. As seen from this plot, the maximum errors are bounded $< 1\%$, and close to the domain bottom boundaries. Averaging across the domain, we have error $e = 5.2 \times 10^{-3}$. This confirms PINN’s ability to capture PDE solutions in higher dimensional domains. Next, we present the PINN loss histories during training. Here, the loss history is constrained by both $\text{MSE}_{\mathcal{P}}$ and MSE_b and decreases slowly with increasing epochs. Both this loss history and the concentration of PDE errors close to the bottom boundaries indicate that the PINN has difficulty adjusting to the boundary condition. Nevertheless, it is successfully able to propagate the boundary information into the domain and capture the interior solutions.

5.6.2 Starfish domain

Finally, we challenge our PINN approach to capture solutions in complex domains. Here we choose a domain resembling the shape of a starfish, shown in the right column of Fig. 2, characterized by the following function

$$r = a(1 + \epsilon \cos(n\theta))$$

where

$$r := \sqrt{x^2 + y^2} \quad \theta := \arctan(y/x)$$

are coordinates in the polar domain, $a = 1, \epsilon = 0.4, n = 5$. Here $r \in [0, r_o]$ and $\theta \in [0, 2\pi]$, with r_o representing the outer domain edge. In this complex shape setting, we choose the following boundary condition,

$$u_b(r_o, \theta) = \sin \theta$$

chosen for its simplicity while retaining (physical) periodicity around the curve. Even with this simple parameterization, we do not have analytical solutions, so we turn to numerical solutions using an integral equations approach.

Here, we train our PINN with $N_{\mathcal{P}} = 2500$, with a uniform distribution of 25 points in the r domain and 100 points in the θ domain. We set $N_b = 100$ sampled on the outermost edge of the domain (shown as blue crosses in the middle row of the figure). We recover the solutions shown in the right column of Fig. 6. Here we once again show the solution generated by the PINN on the first row. Next, we show the spatial distribution of the relative errors (as percentages) obtained upon comparison with the analytical solution. As seen from this plot, the maximum errors are bounded $< 3\%$, and close to the sharp(er) corners of the domain. Averaging across the domain, we have error $e = 3.5 \times 10^{-2}$. This re-confirms PINN’s ability to capture PDE solutions in complex, higher dimensional domains. Next, we present the PINN loss histories during training. Here, the loss history is constrained by boundary losses MSE_b which decreases with increasing epochs. This once again indicates that the PINN has difficulty adjusting to the boundary condition. As a final remark, we note that we can potentially rectify the high boundary losses encountered in both cases shown in this section by penalizing boundary losses more than the PDE loss by setting $\lambda_{\mathcal{P}} < 1$.

Finally, we demonstrate the versatility and robustness of PINNs by solving PDEs on more complex shapes with varying boundary conditions, for a case which is otherwise tedious to solve using conventional numerical algorithms. We chose to parametrize the shape *CS498DL* (chosen for its relevance to the course) and solve the Laplace equation inside the alphanumeric characters. We showcase this solution in Fig. 7. We remark that despite the complexity of this problem we obtain solutions using the same machinery we built up thus far in a straightforward manner, without the need for complicated, explicit treatment at the domain boundaries.



Figure 7: Solution of a Poisson equation generated by PINNs within the shape *CS498DL*. Each alphanumeric character is prescribed with a unique boundary condition resulting in varied, rich solutions seen in the interior.

6 Conclusions

In summary, we successfully implemented and investigated PINN for solving non-linear partial differential equations. We demonstrated PINNs ability to solve equations of increasing complexity, from linear, 1D problems to nonlinear, 2D problems. In particular, we first demonstrated PINN’s utility in solving 1D linear BVPs (Poisson, Helmholtz) with varying boundary conditions and forcing functions, followed by a non-linear BVP (steady viscous Burgers). We then extend this machinery to solve a more complex spatio-temporal non-linear PDE in quasi-2D regime (unsteady viscous Burgers) where we showed that PINNs could successfully capture the expected, non-trivial, non-linear behaviors (shocks and rarefaction waves). We then demonstrated PINN solutions on fully 2D BVPs (Laplace) on regular (rectangular) and arbitrary (starfish) domains. Finally, we conclude by presenting the versatility and robustness of PINN by solving a PDE (Laplace) on a complex CS498 shape and varying boundary conditions, which is otherwise difficult to solve by conventional numerical algorithms.

7 Personal interests

In our first project checkpoint we were asked to “*provide more information on how this topic relates to your backgrounds and any research projects outside of this class*”. This section intends to answer this question.

All team members are currently pursuing doctoral studies in continuum fluid and solid mechanics using predominantly computational techniques. These involve developing algorithms, implementing scalable software and using these to investigate physics for rationally designing new applications. As part of this process, we need to numerically solve multiple coupled PDEs, similar to those demonstrated in this report, for large systems (containing $10^7 - 10^{10}$ degrees of freedom). In this context, we are interested in leveraging ML techniques to either

- rapidly find solutions to such large DOF PDEs
- accelerate present numerical PDE solvers for improved speed
- provide data-driven model closures when the PDE is unknown (such as in turbulence)
- provide black-box control capabilities in devising physics-based applications

8 Contributions

All three members were involved in project conceptualization, implementing the software, running numerical experiments, analyzing and visualizing data and finally preparing the report.

References

- [1] Wikipedia contributors. Partial differential equation — Wikipedia, the free encyclopedia, 2021. [Online; accessed 11-April-2021].
- [2] Alfio Quarteroni and Alberto Valli. *Numerical approximation of partial differential equations*, volume 23. Springer Science & Business Media, 2008.
- [3] M Mitchell Waldrop. More than moore. *Nature*, 530(7589):144–148, 2016.
- [4] Hadi Esmaeilzadeh, Emily Blem, Renee St Amant, Karthikeyan Sankaralingam, and Doug Burger. Dark silicon and the end of multicore scaling. In *2011 38th Annual international symposium on computer architecture (ISCA)*, pages 365–376. IEEE, 2011.
- [5] Yohai Bar-Sinai, Stephan Hoyer, Jason Hickey, and Michael P Brenner. Learning data-driven discretizations for partial differential equations. *Proceedings of the National Academy of Sciences*, 116(31):15344–15349, 2019.
- [6] Maziar Raissi, Paris Perdikaris, and George E Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- [7] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *arXiv preprint arXiv:1912.01703*, 2019.
- [8] Xiaowei Jin, Shengze Cai, Hui Li, and George Em Karniadakis. Nsfnets (navier-stokes flow nets): Physics-informed neural networks for the incompressible navier-stokes equations. *Journal of Computational Physics*, 426:109951.
- [9] Maziar Raissi, Alireza Yazdani, and George Em Karniadakis. Hidden fluid mechanics: Learning velocity and pressure fields from flow visualizations. *Science*, 367(6481):1026–1030, 2020.
- [10] Ameya D Jagtap, Kenji Kawaguchi, and George Em Karniadakis. Adaptive activation functions accelerate convergence in deep and physics-informed neural networks. *Journal of Computational Physics*, 404:109136, 2020.
- [11] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [12] Joe Miroe. Traffic shockwaves, 2017. [Online at youtu.be/6ZC9h8jgSj4; accessed 27-April-2021].